

Program No: 1a. Backpropagation with Feed-Forward Network using Numpy

1. Problem Statement

Feed-forward neural networks cannot correctly solve classification problems like XOR without a learning mechanism. The objective of this problem is to apply backpropagation to reduce prediction error and enable accurate classification.

2. Objective

To design a feed-forward neural network, train it using backpropagation, and evaluate its performance before and after training by analyzing loss values and visualizing error reduction.

3. Dataset

N/A

4. Dataset Analysis

N/A

5. Program

```
import numpy as np
import matplotlib.pyplot as plt
```

```
# Data
```

```
X = np.array([[0,0],
               [0,1],
               [1,0],
               [1,1]])
y = np.array([[0],
               [1],
               [1],
               [0]])
```

```
np.random.seed(42)
```

```
# Initialize weights
```

```
W1 = np.random.randn(2, 3)
b1 = np.zeros((1, 3))
W2 = np.random.randn(3, 1)
b2 = np.zeros((1, 1))
```

```
lr = 0.1
```

```
epochs = 3000
```

```
loss_history = []
```

```
# Activation
```

```
def sigmoid(z):
    return 1 / (1 + np.exp(-z))
```

```
def sigmoid_deriv(a):
    return a * (1 - a)
```

```
# BEFORE BACKPROP (Only forward pass)
z1 = X @ W1 + b1
a1 = sigmoid(z1)
z2 = a1 @ W2 + b2
y_pred_before = sigmoid(z2)
loss_before = np.mean((y - y_pred_before) ** 2)
# TRAINING (Backprop)
for _ in range(epochs):

    # Forward
    z1 = X @ W1 + b1
    a1 = sigmoid(z1)
    z2 = a1 @ W2 + b2
    y_hat = sigmoid(z2)

    loss = np.mean((y - y_hat) ** 2)
    loss_history.append(loss)

    # Backpropagation
    dZ2 = (y_hat - y) * sigmoid_deriv(y_hat)
    dW2 = a1.T @ dZ2
    db2 = np.sum(dZ2, axis=0, keepdims=True)

    dZ1 = (dZ2 @ W2.T) * sigmoid_deriv(a1)
    dW1 = X.T @ dZ1
    db1 = np.sum(dZ1, axis=0, keepdims=True)

    # Update
    W2 -= lr * dW2
    b2 -= lr * db2
    W1 -= lr * dW1
    b1 -= lr * db1
# AFTER BACKPROP
z1 = X @ W1 + b1
a1 = sigmoid(z1)
z2 = a1 @ W2 + b2

y_pred_after = sigmoid(z2)
loss_after = np.mean((y - y_pred_after) ** 2)
print("Loss BEFORE backprop:", loss_before)
print("Loss AFTER backprop :", loss_after)
print("\nPredictions before training:")
print(np.round(y_pred_before, 3))
print("\nPredictions after training:")
print(np.round(y_pred_after, 3))
plt.plot(loss_history)
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Loss Reduction Using Backpropagation")
plt.show()
```

6. Output


```
*** Loss BEFORE backprop: 0.31824520886068175  
Loss AFTER backprop : 0.058651326187686294
```

Predictions before training:

```
[[0.719]  
 [0.807]  
 [0.737]  
 [0.806]]
```

Predictions after training:

```
[[0.164]  
 [0.786]  
 [0.745]  
 [0.311]]
```



7. Performance Analysis (5 factors)

Loss: Reduced from 0.3 to 0.05 as backpropagation epochs increase

Number of Epochs: 3000

8. Result

After training, the loss value decreases significantly, showing that the network has learned effectively. The predicted outputs are very close to the expected XOR target values. This indicates that the model successfully learned the XOR pattern and produces accurate results.

Program No: 1b. Backpropagation with Feed-Forward Network using TensorFlow

1. Problem Statement

A feed-forward neural network produces inaccurate predictions for classification tasks such as XOR. Without a learning mechanism, the network cannot minimize error or improve its performance. This problem focuses on training the network to reduce prediction error using backpropagation.

2. Objective

The objective is to implement a feed-forward neural network and train it using backpropagation, then compare the network's performance before and after training by analysing loss values and visualizing error reduction through a graph.

3. Dataset

N/A

4. Dataset Analysis

N/A

5. Program

```
import tensorflow as tf
X = tf.constant([[0., 0.],
                 [0., 1.],
                 [1., 0.],
                 [1., 1.]], dtype=tf.float32)
y = tf.constant([[0.],
                 [1.],
                 [1.],
                 [0.]], dtype=tf.float32)
tf.random.set_seed(42)
W1 = tf.Variable(tf.random.normal((2, 3)))
b1 = tf.Variable(tf.zeros((1, 3)))
W2 = tf.Variable(tf.random.normal((3, 1)))
b2 = tf.Variable(tf.zeros((1, 1)))
lr = 0.1
epochs = 3000
loss_history = []
def sigmoid(z):
    return tf.nn.sigmoid(z)

# BEFORE BACKPROP (Only forward pass)
z1 = tf.matmul(X, W1) + b1
a1 = sigmoid(z1)
z2 = tf.matmul(a1, W2) + b2
y_pred_before = sigmoid(z2)
loss_before = tf.reduce_mean(tf.square(y - y_pred_before))
# TRAINING (Backpropagation)
for _ in range(epochs):
    with tf.GradientTape() as tape:
        z1 = tf.matmul(X, W1) + b1
        a1 = sigmoid(z1)
```

```
z2 = tf.matmul(a1, W2) + b2
y_hat = sigmoid(z2)
loss = tf.reduce_mean(tf.square(y - y_hat))
grads = tape.gradient(loss, [W1, b1, W2, b2])
```

```
loss_after = tf.reduce_mean(tf.square(y - y_pred_after))
```


```
plt.plot(loss_history)
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Loss Reduction Using Backpropagation (TensorFlow)")
plt.show()
```

6. Output

```
Loss BEFORE backprop: 0.2422114
Loss AFTER backprop : 0.16372998

Predictions before training:
[[0.45 ]
 [0.511]
 [0.502]
 [0.528]]

Predictions after training:
[[0.197]
 [0.538]
 [0.726]
 [0.572]]
```



7. Performance Metrics

Loss: Reduced from 0.24 to 0.16 as backpropagation epochs increase
Number of Epochs: 3000

8. Result

After training, the network achieves a significantly lower loss value, and the predicted outputs closely match the expected XOR targets.

Program No: 2a. Handwritten Digit Recognition using with Keras

1. Problem Statement

Handwritten digit recognition involves classifying image data, which contains spatial and local features that cannot be efficiently learned using simple feed-forward networks. Traditional neural networks fail to exploit spatial patterns present in images

2. Objective

The objective is to design and implement a Convolutional Neural Network using TensorFlow for handwritten digit recognition and to evaluate its performance on unseen test data by measuring classification accuracy.

3. Dataset

MNIST Handwritten Digit Dataset (<https://www.kaggle.com/datasets/hojjatk/mnist-dataset>)

- 60,000 training images
- 10,000 testing images
- Image size: 28×28 pixels
- Grayscale images with pixel values in the range 0–255
- 10 classes representing digits 0–9

4. Dataset Analysis

- Images are normalized to the range [0, 1]
- A channel dimension is added to make the data compatible with CNN input
- Labels are integers from 0 to 9
- The dataset is suitable for multi-class image classification

5. Program

```
import tensorflow as tf
from tensorflow.keras import layers
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0
x_train = x_train[..., None]
x_test = x_test[..., None]
model = tf.keras.Sequential([
    layers.Conv2D(32, (3, 3), activation="relu", input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation="relu"),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(64, activation="relu"),
    layers.Dense(10, activation="softmax")
])
model.compile(
    optimizer="adam",
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)
#Train
model.fit(x_train, y_train, epochs=5, batch_size=64, verbose=1)
#Evaluate on test set
```

```
loss, acc = model.evaluate(x_test, y_test, verbose=0)
print("Test accuracy:", acc)
```

```
#Sample prediction
import numpy as np
x = 0
img = x_test[idx:idx+1]
true_label = y_test[idx]
probs = model.predict(img, verbose=0)[0]
pred_class = np.argmax(probs)
```

```
print("True label:", true_label)
print("Predicted label:", pred_class)
print("Class probabilities:", probs)
```

6. Output

```
Epoch 1/5
938/938 ————— 46s 47ms/step - accuracy: 0.8855 - loss: 0.3906
Epoch 2/5
938/938 ————— 45s 47ms/step - accuracy: 0.9824 - loss: 0.0574
Epoch 3/5
938/938 ————— 85s 50ms/step - accuracy: 0.9886 - loss: 0.0385
Epoch 4/5
938/938 ————— 83s 51ms/step - accuracy: 0.9917 - loss: 0.0278
Epoch 5/5
938/938 ————— 81s 50ms/step - accuracy: 0.9942 - loss: 0.0202
Test accuracy: 0.989899929428101
True label: 7
Predicted label: 7
Class probabilities: [3.3901895e-10 8.7871782e-10 5.8074551e-08 8.7066195e-07 1.8970847e-12
1.2734810e-10 1.4915692e-15 9.9999893e-01 2.0194739e-09 1.8891710e-07]
```

7. Performance Metrics

Batch Size: 938

Number of Epochs: 5

8. Result

The CNN model trained on the MNIST dataset successfully classifies handwritten digits with high accuracy and correctly predicts the digit for a given test image.

Program No: 2b. Handwritten Digit Recognition using with Keras (2 Digits)

1. Problem Statement

Handwritten digit recognition involves classifying image data, which contains spatial and local features that cannot be efficiently learned using simple feed-forward networks. Traditional neural networks fail to exploit spatial patterns present in images

2. Objective

The objective is to design and implement a Convolutional Neural Network using TensorFlow for handwritten digit recognition and to evaluate its performance on unseen test data by measuring classification accuracy.

3. Dataset

MNIST Handwritten Digit Dataset (<https://www.kaggle.com/datasets/hojjatk/mnist-dataset>)

- 60,000 training images
- Image size: 28×28 pixels
- Grayscale images with pixel values in the range 0–255
- 10 classes representing digits 0–9

4. Dataset Analysis

- Images are normalized to the range [0, 1]
- Labels are integers from 0 to 9


5. Program

```
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
import cv2
import numpy as np
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
x_train = x_train / 255.0
x_test = x_test / 255.0
model = models.Sequential([
    layers.Flatten(input_shape=(28, 28)),
    layers.Dense(128, activation='relu'),
    layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
model.fit(x_train, y_train, epochs=5)
test_loss, test_acc = model.evaluate(x_test, y_test)
print("Test Accuracy:", test_acc)
img_path = input("Enter image path (example: two_digits.png): ")
img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
if img is None:
    print("Error: image not found.")
    exit()
img = cv2.resize(img, (56, 28))
img = img / 255.0
digit1 = img[:, :28]
digit1_input = digit1.reshape(1, 28, 28)
```



```
digit2_input = digit2.reshape(1, 28, 28)
pred1 = model.predict(digit1_input)
pred2 = model.predict(digit2_input)
digit1_pred = pred1.argmax()
print("Predicted Number =", digit1_pred, digit2_pred)
plt.figure(figsize=(6,3))
plt.subplot(1,2,1)
plt.imshow(digit1, cmap='gray')
plt.title("Digit 1 = " + str(digit1_pred))
plt.subplot(1,2,2)
plt.show()
```

6. Output



7. Performance Metrics

N/A

8. Result

The CNN model trained on the MNIST dataset successfully classifies 2 handwritten digits with high accuracy and correctly predicts the digit for a given test image.

Program No: 3. Credit Card Fraud Detection

1. Problem Statement

The main difficulty is identifying these fraud cases because they are very few compared to normal transactions. This experiment focuses on using a deep learning model to classify transactions as either normal or fraudulent.

2. Objective

The goal of this experiment is to create a neural network that can detect fraudulent credit card transactions. The model is trained on transaction data and tested using common performance evaluation measures.

3. Dataset

Credit Card Transactions Dataset (<https://www.kaggle.com/datasets/priyamchoksi/credit-card-transactions-dataset>)

4. Dataset Analysis

- The dataset contains highly imbalanced classes, where fraudulent transactions form a very small percentage of the total data.
- All features are numerical, and most features (V1–V28) are anonymized and transformed using Principal Component Analysis (PCA).
- The Class attribute indicates whether a transaction is fraudulent (1) or legitimate (0).

5. Program

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score

data = pd.read_csv("creditcard.csv")

X = data.drop("Class", axis=1)
y = data["Class"]
data = pd.read_csv("creditcard.csv")

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

model = Sequential()
model.add(Dense(16, activation='relu', input_shape=(X_train.shape[1],)))
model.add(Dense(8, activation='relu'))
```

```
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='adam',loss='binary_crossentropy',metrics=['accuracy'])
history = model.fit(X_train, y_train, epochs=10,batch_size=32,validation_split=0.2)
```



```
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Neural Network Loss Curve")
plt.legend()
plt.show()
```

```
y_pred = (model.predict(X_test) > 0.8).astype(int)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.title("Neural Network Accuracy Curve")
plt.legend()
plt.show()
```


```
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

```
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
print("Accuracy :", accuracy)
print("Precision:", precision)
print("Recall :", recall)
print("F1-score :", f1)
```

6. Output



```
Epoch 1/10
5697/5697 - 17s 3ms/step - accuracy: 0.9873 - loss: 0.0605 - val_accuracy: 0.9994 - val_loss: 0.0034
Epoch 2/10
5697/5697 - 16s 3ms/step - accuracy: 0.9993 - loss: 0.0034 - val_accuracy: 0.9993 - val_loss: 0.0031
Epoch 3/10
5697/5697 - 15s 3ms/step - accuracy: 0.9994 - loss: 0.0028 - val_accuracy: 0.9993 - val_loss: 0.0036
Epoch 4/10
5697/5697 - 15s 3ms/step - accuracy: 0.9994 - loss: 0.0028 - val_accuracy: 0.9993 - val_loss: 0.0030
Epoch 5/10
5697/5697 - 15s 3ms/step - accuracy: 0.9995 - loss: 0.0019 - val_accuracy: 0.9993 - val_loss: 0.0030
Epoch 6/10
5697/5697 - 16s 3ms/step - accuracy: 0.9996 - loss: 0.0020 - val_accuracy: 0.9993 - val_loss: 0.0030
Epoch 7/10
5697/5697 - 15s 3ms/step - accuracy: 0.9995 - loss: 0.0023 - val_accuracy: 0.9993 - val_loss: 0.0033
Epoch 8/10
5697/5697 - 15s 3ms/step - accuracy: 0.9995 - loss: 0.0021 - val_accuracy: 0.9993 - val_loss: 0.0033
Epoch 9/10
5697/5697 - 30s 5ms/step - accuracy: 0.9996 - loss: 0.0019 - val_accuracy: 0.9994 - val_loss: 0.0030
Epoch 10/10
5697/5697 - 21s 4ms/step - accuracy: 0.9995 - loss: 0.0019 - val_accuracy: 0.9994 - val_loss: 0.0030
```



7. Performance Metrics

- Accuracy: 0.9993
- Precision: 0.8351
- F1-Score: 0.804

8. Result

The neural network model successfully learned meaningful patterns from the transaction data and was able to accurately differentiate between legitimate and fraudulent transactions. This demonstrates the model's capability to capture complex relationships within the data and effectively support fraud detection tasks.

Program No: 4a. Image Classification

1. Problem Statement

To design and develop a deep learning-based multi-class image classification system that can accurately detect, identify, and categorize multiple object classes within a single image, demonstrating effective performance in real-time visual recognition tasks

2. Objective

To apply a pre-trained YOLOv8 model for object classification in images and analyze its accuracy and performance.

3. Dataset

N/A


4. Dataset Analysis

N/A

5. Program

```
!pip install ultralytics
import cv2
from ultralytics import YOLO
from google.colab.patches import cv2_imshow
yolo = YOLO('yolov8n.pt')
image_path = input("Enter the path to your image: ") # The image path
results = yolo(image_path)
for r in results:
    im_array = r.plot() # plot a BGR numpy array of predictions
    # OpenCV loads images as BGR, but cv2_imshow expects RGB for correct color display.
    im_rgb = cv2.cvtColor(im_array, cv2.COLOR_BGR2RGB)
    cv2_imshow(im_rgb)
```

6. Output



7. Performance Metrics

- Accuracy: 0.83
- Speed: 14.4ms (preprocess), 325.1ms (inference)

8. Result

Using YOLOv8, the image classification system was implemented successfully, and the model identified the object as a bus with high confidence, highlighting its real-time classification capability.

Program No: 4b. Multi-class Image Classification

1. Problem Statement

To design and develop a deep learning-based multi-class image classification system capable of identifying and classifying an image into one of several predefined object categories.

2. Objective

To classify an input image into a single object class using a pre-trained YOLOv8 model and analyze its classification performance.

3. Dataset

N/A


4. Dataset Analysis

N/A

5. Program

```
!pip install ultralytics
import cv2
from ultralytics import YOLO
from google.colab.patches import cv2_imshow
yolo = YOLO('yolov8n.pt')
# Input image path
image_path = input("Enter the path to your image: ")
# Perform object detection
results = yolo(image_path)
# Display results
for r in results:
    im_array = r.plot() # Plot predictions (BGR format)
    # Convert BGR to RGB for correct display
    im_rgb = cv2.cvtColor(im_array, cv2.COLOR_BGR2RGB)
    cv2_imshow(im_rgb)
```

6. Output



7. Performance Metrics

- Accuracy: 0.80, 0.84
- Speed: 7.2ms (preprocess), 226.4ms (inference)

8. Result

The multi-class image classification system was successfully implemented using YOLOv8.

Program No: 5. Student Attendance System Using Face Recognition

1. Problem Statement

To design a deep learning–based face recognition system for automatic attendance marking by identifying and verifying student faces from images.

2. Objective

To create an automated student attendance system based on face recognition and analyze how accurately it identifies registered students.

3. Dataset

5 images selected from the Human Faces Dataset

(<https://www.kaggle.com/datasets/kaustubhdhote/human-faces-dataset>)

4. Dataset Analysis

Total folder consists of 4360 AI Generated Images & 5000 real images. 5 images were chosen from the real images sub-folder.

5. Program

```
!pip install face_recognition opencv-python
```

```
import face_recognition
```

```
import cv2
```

```
import os
```

```
import numpy as np
```

```
# Lists to store known faces
```

```
known_face_encodings = []
```

```
known_face_names = []
```

```
# Path where all images are stored
```

```
dataset_path = "students_dataset"
```

```
# Load each image from the folder
```

```
for image_name in os.listdir(dataset_path):
```

```
    image_path = os.path.join(dataset_path, image_name)
```

```
    # Load image
```

```
    image = face_recognition.load_image_file(image_path)
```

```
    # Encode face
```

```
    encoding = face_recognition.face_encodings(image)[0]
```

```
    # Extract student name from filename
```

```
    student_name = os.path.splitext(image_name)[0]
```

```
    known_face_encodings.append(encoding)
```

```
    known_face_names.append(student_name)
```

```
# Load test image
```

```
test_image_path = input("Enter the path to test image: ")
test_image = face_recognition.load_image_file(test_image_path)

# Detect faces
face_locations = face_recognition.face_locations(test_image)
face_encodings = face_recognition.face_encodings(test_image, face_locations)

# Convert image for display
image_bgr = cv2.cvtColor(test_image, cv2.COLOR_RGB2BGR)


# Recognize faces
for (top, right, bottom, left), face_encoding in zip(face_locations, face_encodings):
    matches = face_recognition.compare_faces(known_face_encodings, face_encoding)
    name = "Unknown"

    if True in matches:
        first_match_index = matches.index(True)
        name = known_face_names[first_match_index]

    cv2.rectangle(image_bgr, (left, top), (right, bottom), (0, 255, 0), 2)
    cv2.putText(image_bgr, name, (left, top - 10),
                cv2.FONT_HERSHEY_SIMPLEX, 0.9, (0, 255, 0), 2)

cv2.imshow("Attendance System", image_bgr)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

6. Output



7. Performance Metrics

N/A

8. Result

The student attendance system using face recognition was successfully implemented. The system correctly identified registered students from facial images and automatically recorded their attendance.